

Jigsaw: Memory-Bounded Exact Subgraph Matching on Million-Node Attributed Graphs

Swastik Mohapatra
PES University
Bengaluru, India
swastik9895@gmail.com

ABSTRACT

Exact subgraph matching enumerates every embedding of a small attributed query graph in a large data graph, and it underpins graph databases, knowledge-graph search, and program analysis. Exact solvers are memory-bound: run over the whole graph, the Glasgow Subgraph Solver enumerates all 15 test queries on CoraFull (19.8K nodes) but none once the graph reaches OGBN-Arxiv scale (169K nodes), and at a million nodes holding the graph together with the solver’s candidate domains resident becomes impractical. We present **Jigsaw**, a retrieval-constrained matcher that keeps exact matching feasible under a fixed memory budget. Jigsaw partitions the data graph offline, retrieves a small set of candidate partitions per query with a coverage-trained encoder, expands them by one-hop boundary overlap, prunes by typed signature and exact label, and invokes the exact solver only inside the retrieved candidate. The guarantee is precise: every returned mapping is an exact match within the retrieved candidate under a configured discrete label semantics, verified by the solver alone, and the memory budget K trades recall for footprint but never soundness; a bounded no-match result is not a certificate of global non-existence. On OGBN-MAG (1.9M nodes) Jigsaw solves **88.6%** of 1,800 planted positives while a two-pass streaming implementation holds **2.4 GB** instead of **10.2 GB** for whole-graph residence (4.2 $\times$ less). Because the value of learning depends on label selectivity, Jigsaw routes retrieval by an offline selectivity statistic: a classical label index dominates when labels are near-unique (96.7% on MAG) and the coverage-trained encoder dominates under realistic class labels. We report both rather than a single winner, turning retrieval into a practical execution layer for exact verification rather than a neural replacement for it.

PVLDB Reference Format:

Swastik Mohapatra. Jigsaw: Memory-Bounded Exact Subgraph Matching on Million-Node Attributed Graphs. PVLDB, 20(1): XXX-XXX, 2027.
doi:XX.XX/XXX.XX

PVLDB Artifact Availability:

The source code, data, and/or other artifacts have been made available at <https://github.com/DoggoWoofWoof/Jigsaw>.

This work is licensed under the Creative Commons BY-NC-ND 4.0 International License. Visit <https://creativecommons.org/licenses/by-nc-nd/4.0/> to view a copy of this license. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org. Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.
Proceedings of the VLDB Endowment, Vol. 20, No. 1 ISSN 2150-8097.
doi:XX.XX/XXX.XX

1 INTRODUCTION

Subgraph isomorphism is a basic tool for motif discovery, scientific graph analysis, fraud search, and security auditing. The exact decision problem is NP-complete [9]; modern constraint solvers such as Glasgow [1] are strong, but large candidate domains and whole-graph residence still make them difficult to serve at scale. Learned retrieval can identify promising graph regions, but it cannot certify a match. Our question is whether retrieval can reduce the target enough for exact verification to fit a practical memory budget. Jigsaw uses the neural model only to choose where to search and leaves acceptance to Glasgow.

We call this design *retrieval-constrained verification*. Jigsaw partitions a large graph, embeds queries and partitions in a shared space, retrieves a bounded set of partitions, stitches their overlap neighborhoods, prunes by typed structural signatures and exact labels, and then calls the verifier. If the candidate contains a valid embedding, Glasgow can certify it exactly. Missing a required region limits global completeness but cannot create a false match. FullCov trains the retriever against that specific failure event by requiring coverage of every partition touched by the planted match.

We evaluate the full path across CoraFull, OGBN-Arxiv, and MAG using two locked seeds, three sizes, six positive and two negative families, and six retrieval policies. At matched half-partition budgets the full cascade solves 94.2/92.8/88.6% of planted positives on CoraFull/Arxiv/MAG (Table 2). Two operators put it there: coverage-aware training improves the fixed-budget retrieval operating points (Table 1), and one-hop boundary overlap is the load-bearing recall step (Table 9). Examining every partition reaches 100% on CoraFull/Arxiv but is a recall ceiling, not a scalable policy.

Contributions.

- We turn exact matching into a **memory-bounded service**: partition-wise streaming verifies OGBN-MAG (1.9M nodes) within a 2.4 GB budget versus 10.2 GB for whole-graph residence (4.2 $\times$ less) and exposes a tunable memory-latency frontier that resident and prior learned-exact matchers do not. Under this budget, retrieval-constrained verification recovers 88.6% of the 1,800 planted MAG positives, where a direct whole-graph Glasgow call solves 0/15 of the K-hop timeout control (Table 3); the same substitution turns Arxiv’s direct 0/15 into 15/15.
- We formulate the pipeline as coverage-conditioned exact verification, separating verifier soundness under the configured label semantics (always guaranteed) from retrieval completeness (a tunable budget knob), and prove the pruning stage is coverage-lossless.

- We introduce **FullCov**, which targets the weakest required partitions: Conditional Value-at-Risk (CVaR) averages the largest tail of their losses, and a top- K barrier penalizes those below the retrieval cutoff.
- We provide a two-seed, three-dataset benchmark with positive and negative query families, exact-verification metrics, and an audited evaluation, measuring each candidate-shrinkage stage to show that boundary overlap is essential for recall while query-derived pruning preserves every covered match.

2 PROBLEM AND GUARANTEE

Let $G = (V, E, X)$ be a large attributed data graph and $Q = (V_Q, E_Q, X_Q)$ a query graph. The exact task asks for embeddings of Q into G that preserve edges, node types, and a discrete vertex-label map σ under the solver’s configured semantics. Whole-graph exact search is memory-bound: the verifier must hold G and its candidate domains resident, which becomes infeasible as G reaches millions of nodes (Section 5). Jigsaw instead constructs a candidate subgraph $G_C \subseteq G$ and runs the exact verifier only on G_C . Three claims are commonly conflated; we state precisely which ones Jigsaw guarantees.

What Jigsaw guarantees

Fix a data graph G , query Q , label semantics σ , and memory budget K ; Jigsaw returns a set M of embeddings of Q into G . **(i) Soundness.** Every $m \in M$ is an exact embedding under σ , certified by the Glasgow verifier on G_C alone; no learned component can accept a match. The default σ is an attribute hash, so acceptance is exact *under* σ by construction; raw-attribute exactness is stronger, holds empirically on this benchmark (audited below), and is guaranteed by construction only with the optional raw-equality guard enabled. **(ii) Candidate-local completeness.** M contains every embedding whose image lies inside the retrieved candidate G_C ; matching-invariant pruning removes no valid embedding (Proposition 5.1). **(iii) Budget monotonicity.** Enlarging K can only grow G_C , hence M , and never admits a spurious match; at $K = |\mathcal{P}|$ every partition is admitted and the run is complete under σ whenever the verifier finishes. Jigsaw does *not* certify global non-existence: an empty M means “no match inside G_C ,” not “no match in G .” Recall of true matches is an empirical quantity we measure, not a worst-case bound.

Coverage metric. Retrieval quality is governed by whether the candidate covers the match. Let $\mathcal{T}(Q)$ be the coarse partitions touched by a planted embedding and $R_K(Q)$ the retrieved top- K partition set. We define

$$\text{FullCov}@K(Q) = 1 [\mathcal{T}(Q) \subseteq R_K(Q)]. \quad (1)$$

Partition recall@ K is $|\mathcal{T}(Q) \cap R_K(Q)|/|\mathcal{T}(Q)|$; FullCov@ K equals one only when recall is complete. One-hop boundary overlap expands R_K before pruning, so FullCov is sufficient but not necessary: overlap can recover a required boundary node even when its home partition is not in the top K .

Budget as a soundness-preserving completeness dial. Soundness never depends on K : any returned match is a true match at every

budget, so K is a pure completeness control. At $K = |\mathcal{P}|$ (FilterAll) retrieval covers the whole partition set and matching-invariant pruning preserves every valid embedding, so the run is complete under σ whenever the verifier finishes. As K shrinks, the solved-by-budget curves of Section 5 measure realized completeness under the timeout. A caller needing certainty pays for all partitions and sufficient verifier time; one needing speed accepts a quantified miss/timeout rate.

Label semantics and why acceptance is exact. The verifier consumes discrete labels $\sigma(v)$; by default we set $\sigma(v) = \text{md5}(x_v) \bmod 10^6$, a hash of the full node attributes. Discretization is used only to shrink the retrieval and pruning domain, and because pruning is matching-invariant (Proposition 5.1) a hash collision can only enlarge the candidate the verifier inspects, never remove a valid embedding. Two exactness notions must be separated. Under σ -semantics the Glasgow certificate is exact by definition. Raw-attribute exactness is stronger, and the modular hash *does* collide: across MAG’s 742,524 distinct node representations, $\text{md5} \bmod 10^6$ yields only 524,328 distinct codes, so 52.3% of representations share a bucket and 34.4% of nodes fall in a colliding bucket. A σ -certified mapping onto a collided node would be σ -exact but not necessarily raw-exact. We address this in two ways: (a) we audit every accepted mapping (1,595 solved positives, 87,536 mapped node pairs) and find zero raw-attribute mismatches on this benchmark—an empirical result on the tested queries, not a worst-case bound; and (b) a drop-in raw-attribute equality guard on each returned mapping restores raw-exactness by construction for any σ . (Separately, a dataset-integrity check confirms all 1,800 planted witnesses, 101,951 pairs.) Section 5 additionally reports a coarser 39-class label regime, which stresses selectivity and motivates the retriever portfolio.

3 JIGSAW METHOD

Figure 1 gives the dataflow. Offline, METIS partitions G into fine and coarse blocks; the shared encoder embeds every coarse block once and FAISS indexes the normalized vectors. Online, Jigsaw embeds the query, retrieves coarse blocks, stitches boundaries, prunes, and verifies (Algorithm 1).

Hierarchical localization and overlap cascade. Fine blocks give local structure; coarse blocks are the retrieval units. For each query Jigsaw computes $z_Q = f_\theta(Q)$ and ranks all coarse partitions by cosine similarity to the cached partition embeddings, retrieving the top K . True query neighborhoods often cross hard partition boundaries, so Jigsaw stitches each retrieved partition with its one-hop boundary nodes from adjacent partitions – not the adjacent partitions in full – raising recall before verification while still avoiding whole-graph search.

Pruning and verification. The stitched region is filtered by typed attribute signatures (node type plus hashed feature fingerprints) and exact labels before Glasgow verification; connected components become independent verifier calls. Both token sets are derived exclusively from the submitted query payload – planted target-node IDs are evaluation metadata and never enter candidate construction or serving-time pruning. Only the exact verifier returns a match. As a running example, suppose a planted query touches blocks p_4 and p_7 : if both rank in the top K , FullCov holds; if only p_4 ranks but the




Figure 1: Jigsaw pipeline. Offline work partitions and embeds the target; online work embeds a query, retrieves coarse partitions, adds one-hop boundary overlap, prunes with query-derived invariants, and invokes Glasgow. The learned component ranks regions, but only Glasgow accepts a match. FullCov is the condition under which the candidate contains every partition the planted match needs.

Algorithm 1 Jigsaw online serving (one query)

Require: target G partitioned into coarse blocks $\mathcal{P}$ with cached embeddings $\{p_j\}$ and FAISS index; encoder f_θ ; query Q ; budget K ; label map σ

Ensure: set M of exact embeddings of Q inside the retrieved candidate

- 1: $z_Q \leftarrow f_\theta(Q)$ ▷ embed query
 - 2: $R \leftarrow \text{TopK}(\cos(z_Q, \{p_j\}), K)$ ▷ retrieve coarse blocks
 - 3: $G_C \leftarrow \bigcup_{p \in R} p \cup \text{ONEHOPBOUNDARY}(R)$ ▷ stitch boundary overlap
 - 4: $G'_C \leftarrow \{v \in G_C : \sigma(v) \in \sigma(V_Q) \wedge \text{sig}(v) \in \text{sig}(V_Q)\}$ ▷ lossless prune (Prop. 5.1)
 - 5: $M \leftarrow \emptyset$
 - 6: **for** each connected component C of G'_C **do**
 - 7: $M \leftarrow M \cup \text{GLASGOW}(Q, C)$ ▷ exact verify; sole accepting step
 - 8: **end for**
 - 9: **return** M
-

required nodes of p_7 lie on its boundary, overlap still completes the candidate; if an interior node of p_7 is never admitted, the bounded run may miss the match but cannot return an invalid one.

FullCov objective. Standard contrastive training [43] rewards high average similarity to positives, which can still leave one required partition ranked too low; Jigsaw optimizes the bottleneck. For query i , let P_i be its required partitions and $u_{ij} = z_i^\top p_j / \tau_{\text{cov}}$ the temperature-scaled similarity to partition j (τ_{cov} a fixed temperature). Each positive $j \in P_i$ receives

$$c_{ij} = -\log \frac{e^{u_{ij}}}{\sum_m e^{u_{im}}}, \quad h_i = \max_{n \notin P_i} u_{in}, \quad (2)$$

$$\ell_{ij} = c_{ij} + 0.25 \text{softplus}(h_i - u_{ij}).$$

For tail fraction ρ , CVaR_ρ [44] averages the largest $\lceil \rho |P_i| \rceil$ losses, so easy positives cannot hide the weakest one; a top- K barrier $b_{ij} = \text{softplus}(t_i(K) - u_{ij})$ pushes borderline positives above the cutoff $t_i(K)$, defined as the similarity of the K -th ranked partition for query i (the retrieval boundary a positive must clear). The coverage

loss is

$$\mathcal{L}_{\text{cov}} = \text{CVaR}_\rho(\{\ell_{ij}\}) + 0.35 \text{CVaR}_\rho(\{b_{ij}\}). \quad (3)$$

The complete objective adds level-specific contrastive terms,

$$\mathcal{L} = 0.2 \mathcal{L}_{\text{fine-NCE}} + 0.8 \mathcal{L}_{\text{coarse-NCE}} + \mathcal{L}_{\text{cov}}, \quad (4)$$

where the coarse term dominates because coarse partitions are the retrieval unit the verifier consumes, the fine term shapes the same space at finer granularity, and node-alignment terms are switched off (weight 0) because node embeddings are unused at query time. The coefficients (the 0.25 margin, the 0.35 barrier, and the 0.2/0.8 fine/coarse split) are held fixed across datasets; only the budget K and tail fraction ρ vary (Arxiv $K=20, \rho=0.25$; MAG $K=50, \rho=0.5$, reflecting its larger footprints). Optimizing the tail rather than the average follows from the metric: FullCov is a conjunction over required partitions, so a query fails as soon as its worst-ranked required partition leaves the budget—a min over required partitions in the hard limit, of which CVaR_ρ is a smooth relaxation. That worst case is exactly the failure exact verification cannot tolerate.

3.1 Encoder architecture

The encoder is shared between query and partition graphs: node features are projected to 256 dimensions, passed through six residual message-passing blocks with LayerNorm, and read out by multi-layer mean/max/sum pooling into an L_2 -normalized 128-d vector (Figure 2). This matters because the task is region localization, not node alignment: the model must place a query near *every* partition needed for verification, including the farthest true one. For CoraFull and Arxiv the block is GraphSAGE [3]. OGBN-MAG is heterogeneous (paper/author/institution/field nodes, typed relations); we flatten to a single index but retain node and relation types through a six-layer RGCN [4] (per-relation bases over symmetrized relations). Since only paper nodes carry features, others receive node-type one-hot and per-relation-degree features, so the retriever reasons over typed structure. The residual connections and LayerNorm are not incidental: Section 5.4 shows they are what keeps a six-layer encoder from oversmoothing. The architecture does not make the retriever a solver; it proposes a candidate boundary, and once G_C is built Glasgow gives exact answers for G_C .




Figure 2: Shared query/partition encoder. One network embeds both queries and cached partitions into a common space: six residual GraphSAGE blocks for CoraFull/Arxiv, a relation-aware RGCN for MAG, pooled to an L_2 -normalized vector and scored by cosine similarity.

4 EXPERIMENTAL PROTOCOL

Datasets. We evaluate CoraFull [13] (small enough for exact timing sanity checks), OGBN-Arxiv [8] (169,343 nodes, the medium-scale benchmark), and the full heterogeneous OGBN-MAG (1,939,743 nodes across four node types joined by typed relations), which we model relation-aware rather than collapsing to attribute-only structure, stressing both scale and heterogeneity.

Queries and policies. All positive families are generated as connected planted subgraphs. The matrix uses two query seeds, sizes 20/50/100, and 50 queries per type/size/seed. Positive families are single, multi-fine, K-hop, degree-K-hop, random-walk, and multi-coarse; negatives perturb labels or structure. Retrieval policies are learned Jigsaw, MeanFeat, Mean-RRF (reciprocal-rank fusion of learned and MeanFeat rankings, a configuration of our retriever), TopoFeat, Random, and exhaustive FilterAll; on MAG we additionally run FeatureIndex, a classical label-coverage index. Model selection is protocol-disjoint: the deployed MAG encoder is the checkpoint chosen by validation FullCov on seeds 31415/27182, while all benchmark queries use seeds 20260607/20260608, so no benchmark query influenced training, weighting, or checkpoint choice.

Partition and overlap scale. On MAG a coarse partition has median 969 nodes and borders many partitions, but one-hop overlap admits only boundary nodes (median 4,252), not whole neighboring partitions; the loose union bound is never materialized, and pruning leaves much smaller verifier candidates (Section 5).

Metrics and hygiene. We report FullCov@ K /recall for retrieval and, for the full path, exact-solve rate, audited no-match accuracy, false positives, timeouts, candidate size, and wall-clock components. Each dataset has 2,400 unique queries (1,800 positives, 600 negatives); across three datasets this is 7,200 unique queries. Timings are per-query wall-clock on 8-vCPU Google Cloud instances with a 30 s verifier budget. Encoders, query seeds, hierarchies, and manifests will be released.

5 RESULTS

We first isolate the retriever (Section 5.1), then report the end-to-end system across three datasets and the label-selectivity regime

Table 1: Matched Arxiv retrieval objective ablation. Fixed retrieval budgets on $n=400$ locked K-hop queries (200 per query seed); each row averages the validation-selected checkpoint from two independent training replicates.

Objective	FullCov@20	FullCov@50	FullCov@100	Mean worst rank
Contrastive control	21.7%	39.0%	66.0%	74.02
FullCov objective	26.0%	48.8%	82.0%	58.97
McNemar p	4.6×10^{-3}	9.8×10^{-6}	4.0×10^{-10}	–

(Section 5.2), the headline memory result (Section 5.3), the encoder-depth diagnostics reviewers of learned components ask for (Section 5.4), the candidate-shrinkage cascade (Section 5.5), query families, operator necessity, and soundness.

5.1 FullCov Training Improves Retrieval

Table 1 gives the matched-budget Arxiv retrieval result. All rows use the same locked K-hop query seeds and 9,000 optimizer steps. The FullCov objective improves the fixed-budget operating points most relevant to bounded exact verification.

The locked K-hop suite has zero queries impossible at $K=20, 50, 100$; average true coarse footprint is 6.37 and maximum footprint is 17. The objective therefore improves ranking by explicitly optimizing the weakest required partition. Because the control and FullCov-aligned models share query seeds and training budget, the gain comes from the objective rather than easier queries or longer optimization. The largest improvement appears at FullCov@100, where the downstream verifier has enough room to benefit from better ordering without searching the whole graph – exactly the cascade’s operating point.

5.2 Three-Dataset Main Benchmark

Table 2 reports exact planted-match recovery, correct no-match behavior, false positives, solver timeouts, candidate nodes, cascade time, and verifier time. Cascade time sums candidate construction and exact verification and excludes retrieval/ranking.

At the half-partition budget (Table 2) Jigsaw solves 94.2/92.8/88.6% on CoraFull/Arxiv/MAG. The 100% CoraFull/Arxiv entries are the exhaustive ceiling, not a retrieval result; Mean-RRF is competitive there, and FeatureIndex is strongest on MAG under near-unique labels. Matched CoraFull/Arxiv costs are measured for every policy: Jigsaw uses mean terminal candidates 458/77, cascade times 0.05/0.19 s, and verifier times 11.8/18.3 ms. Size 100 is harder (CoraFull 89.0% vs. 96.8/96.8% at 20/50; Arxiv 90.3% vs. 93.8/94.2%), driven by CoraFull random-walk (75.3%) and Arxiv multi-coarse (78.7%). On MAG, recovery declines from 93.2% at size 20 to 90.7% at 50 and 82.0% at 100, driven by random-walk, while contained multi-fine remains perfect at every size and multi-coarse remains difficult. Jigsaw sends far fewer candidate nodes than FilterAll (138K vs 444K) at lower verifier time (82 vs 289 ms), while FilterAll remains the recall ceiling at 98.4%.

The scaling break. The central systems result is that whole-graph exact matching does not scale: direct CP-based Glasgow solves all 15 CoraFull K-hop queries but **0 of 15** on Arxiv and MAG (Table 3, Figure 3), whereas the Jigsaw cascade remains feasible by handing

Table 2: Main-benchmark metrics. Pos. uses matched half-partition budgets $K=10/100/1000$; FilterAll uses the all-partition ceiling. Neg. is audited no-match accuracy; FP counts negative false positives, while TO counts positive-query solver timeouts (FeatureIndex was run on positives only). Candidate and timing columns use the same endpoint as Pos. Cascade time is candidate construction plus exact verification and excludes retrieval/ranking. Bold marks Jigsaw, the proposed method (not necessarily the per-column maximum); FilterAll is the unbounded exhaustive ceiling and FeatureIndex applies only under near-unique labels.

Dataset	Method	Pos. %	Neg. %	FP	TO	Cand. nodes	Cascade s	Solver ms
CoraFull	Jigsaw	94.2	100.0	0	0	458	0.05	11.8
CoraFull	Mean-RRF	93.8	100.0	0	0	485	0.05	11.8
CoraFull	MeanFeat	87.7	100.0	0	0	1.13K	0.06	10.5
CoraFull	TopoFeat	31.9	100.0	0	0	7.60K	0.17	3.6
CoraFull	Random	32.7	100.0	0	0	7.54K	0.16	3.8
CoraFull	FilterAll [‡]	100.0	100.0	0	0	59	0.06	12.7
Arxiv	Jigsaw	92.8	100.0	0	0	77	0.19	18.3
Arxiv	Mean-RRF	94.0	100.0	0	0	75	0.19	18.8
Arxiv	MeanFeat	89.7	100.0	0	0	91	0.22	18.4
Arxiv	TopoFeat	33.3	100.0	0	0	233	0.46	6.0
Arxiv	Random	39.2	100.0	0	0	219	0.47	7.0
Arxiv	FilterAll [‡]	100.0	100.0	0	0	57	0.35	19.1
MAG	Jigsaw	88.6	99.5	1	25	138.1K	6.97	81.9
MAG	Mean-RRF	85.4	99.5	1	17	137.1K	5.38	67.9
MAG	FeatureIndex [§]	96.7	—	—	24	59.4K	2.39	118.8
MAG	MeanFeat	64.6	99.5	1	12	212.1K	8.02	27.5
MAG	TopoFeat	37.1	99.8	1	1	235.6K	8.07	5.2
MAG	Random	51.4	99.7	1	9	282.1K	9.54	19.9
MAG	FilterAll [‡]	98.4	99.3	1	33	443.8K	5.85	288.6

[‡]Exhaustive all-partition ceiling (not a bounded-memory policy). [§]Classical label-coverage index; positives only.




Figure 3: Jigsaw makes the exact solver’s domain tractable. Direct full-graph Glasgow solves CoraFull but none of the Arxiv/MAG K-hop samples; Jigsaw stays feasible by verifying only the retrieved-and-pruned candidate.

the verifier a small candidate. The retriever therefore sits on a recall-cost frontier: FeatureIndex is strongest under near-unique MAG labels, FilterAll is the exhaustive ceiling, and the learned retriever is a useful memory-bounded point, especially once labels coarsen.

Label-selectivity crossover. Near-unique md5 labels favor exact-label methods: a classical label-coverage index (**FeatureIndex**, in the GraphGrep [41]/gIndex [42] tradition) is the strongest bounded retriever here (96.7%; median worst-true-partition rank 104 of 2,000) versus the learned GNN (88.6%; rank ~ 900). At coarse 39-class labels this ranking order *inverts*: the learned model has better worst-true-partition ranks in each family, while complete solve runs are statistically tied (FeatureIndex 63.4%; learned variants 63–64%).

Table 3: Scaling and latency: direct full-graph solver vs. Jigsaw cascade. CoraFull/Arxiv use the same 15 K-hop queries; Jigsaw is evaluated at matched half-partition budgets ($K=10/100$). The MAG direct cell is its 15-query timeout control; *the MAG Jigsaw cell is the 1,800-positive main-benchmark mean at $K=1,000$.

Dataset	Nodes	Direct full-graph solver		Jigsaw cascade		Cand. nodes
		Solved	Med. time	Solved	Med. time	
CoraFull	19.8K	15/15	2.74 s	15/15	0.03 s	50
Arxiv	169K	0/15 [†]	≈ 49 s	15/15	0.11 s	50
MAG	1.9M	0/15	timeout	88.6%	6.97 s*	—

This isolates the regime in which learning contributes: as labels stop localizing, the GNN supplies the better partition ranking. The crossover generalizes on CoraFull and Arxiv, where the winner flips with an offline label-selectivity statistic (median partitions/label 1 at feature labels versus 12/124.5 at class labels), so Jigsaw selects its bounded retriever by label regime (Table 4). Jigsaw is therefore *retriever-agnostic*: it routes to a classical index under near-unique labels and to the learned ranker when labels cease to localize.

Statistical significance. Table 5 reports bootstrap confidence intervals and paired exact-McNemar tests for MAG positives. Jigsaw significantly beats every learned/feature retriever except FilterAll,

Table 4: Cross-dataset label-selectivity crossover: overall median worst-true-partition rank (lower is better; retrieval-only, $n=540$ per dataset). Near-unique feature labels favor the classical index; realistic class labels favor the learned retriever.

Dataset (parts)	FeatureIndex		Learned (any)	Parts/label feat / class
	feature	class		
CoraFull (20)	2	7	5	1 / 12
Arxiv (200)	4	104	44	1 / 124.5

Table 5: MAG solve-rate significance. Bootstrap 95% CI over $n=1800$ positive queries, and paired exact-McNemar of Jigsaw vs each policy (wins = queries Jigsaw solves and the other does not). The header solve rate and CI are the deployed encoder (88.6%, 1,595/1,800); in the paired column Mean-RRF is compared against that deployed run, while MeanFeat/Random/Topo/FilterAll are compared against the encoder-independent Jigsaw run (86.0%, 1,548/1,800).

Method	Solve %	95% CI	vs Jigsaw (W/L, p)
Jigsaw	88.6	[87.1, 90.0]	–
Mean-RRF	85.4	[83.7, 87.0]	92 / 35, $p=4.4\times 10^{-7}$
MeanFeat	64.6	[62.3, 66.8]	437 / 51, $p<10^{-76}$
Random	51.4	[49.1, 53.8]	689 / 67, $p<10^{-130}$
Topo	37.1	[34.9, 39.5]	912 / 32, $p<10^{-224}$
FilterAll	98.4	[97.8, 98.9]	2 / 225, $p=2.4\times 10^{-64}$

the exhaustive ceiling. These are component comparisons; the system contribution is exact matching without whole-graph residence and with fewer candidates than FilterAll.

Coarse-label stress on resident pipelines. The same low-selectivity transition breaks resident exact matchers. With MAG labels coarsened to 32 deterministic buckets, resident CFL-Match [19]/DP-iso [20]/GQL [21] complete 12/27 executions (8/9 at size 20; 2/9 at each of sizes 50 and 100; random-walk 3/9), with candidate tables reaching 2.4M entries and 15 calls hitting the 120 s limit. Under near-unique labels the same matchers complete 27/27: a resident exact solver dominates when selective labels collapse its candidate tables, while under weak labels whole-graph candidate construction is what fails, and Jigsaw’s bounded candidate supplies the systems advantage.

5.3 Memory-Bounded Serving

The headline systems benefit is memory. The cascade streams the graph in a two-pass sweep (nodes, then edges) behind an LRU partition cache, so peak resident memory is bounded by the cache rather than by $|G|$. On a 54-query MAG grid (6 families $\times$ 3 sizes $\times$ 3 queries) the conservative primary serve peaks at **2.4 GB** versus **10.2 GB** for whole-graph residence (4.2 $\times$ smaller), while recovering 46/54 matches; an optimized cache sweep lowers the eight-partition point to 1.9 GB (5.4 $\times$). Memory becomes a tunable

Table 6: Depth diagnostics on deployed encoders (no retrain). Partition-summary effective rank and mean pairwise distance from input to layer 6 (higher is better; these are what retrieval discriminates), and the layer-6 node cosine with residual connections disabled ($\rightarrow 1$ is textbook collapse). No dataset oversmooths; on GraphSAGE the residual is the anti-collapse mechanism, while the MAG RGCN is robust without it.

Dataset (enc.)	Part. rank		Part. dist	Resid-off
	input	L6	input $\rightarrow$ L6	L6 node cos
CoraFull (SAGE)	15.2	14.8	0.054 $\rightarrow$ 0.112	0.965 (collapse)
Arxiv (SAGE)	65.0	87.5	0.026 $\rightarrow$ 0.163	0.951 (collapse)
MAG (RGCN)	74.5	118.9	0.062 $\rightarrow$ 0.088	0.586 (robust)

resource: as the cache grows from 8 to 256 partitions, MAG p50 latency falls 5.9 $\rightarrow$ 4.8 s (below the 8.4 s cold load) while the mean holds 17.9 $\rightarrow$ 17.6 s because of a hard-query I/O tail (Figure 4c). Cold-load end-to-end latency is competitive with resident classical matchers – 2.35/3.0/1.2 $\times$ faster on CoraFull/Arxiv/MAG because Jigsaw avoids the full-graph load – so bounding residence does not cost latency for cold or one-shot queries. A resident matcher, by contrast, amortizes its one-time load and then answers in ~ 1 ms, so on long repeated-query streams its steady state is faster; Jigsaw’s decisive advantage is memory, not latency. An all-or-nothing resident matcher cannot expose this frontier: it must hold the whole graph or fail, whereas Jigsaw lets the operator pick a memory $\leftrightarrow$ latency point.

5.4 Encoder Depth, Oversmoothing, and Residual Necessity

Because Jigsaw depends on a six-layer message-passing encoder, a natural concern is oversmoothing: whether depth collapses node and partition representations. We probe the *deployed* weights (no retrain) and measure, at every layer, node-level mean absolute deviation (MAD = $1 - \text{mean off-diagonal cosine}$; $\rightarrow 0$ is collapse) and the retrieval-relevant quantity – effective rank and mean pairwise distance of the mean-pooled *partition* summaries FAISS actually discriminates (Table 6, Figure 5).

Three findings result. **(1) No oversmoothing at six layers on any dataset.** Partition summaries stay fully discriminable and in fact separate *more* with depth: CoraFull partition distance rises 0.054 $\rightarrow$ 0.112, and MAG partition rank rises 74.5 $\rightarrow$ 118.9 with distance 0.062 $\rightarrow$ 0.088. **(2) On Arxiv, depth actively helps** – every metric improves monotonically (node MAD 0.15 $\rightarrow$ 0.59, partition distance 0.026 $\rightarrow$ 0.163, +240%), so six layers are justified rather than harmful. **(3) Residual connections are the mechanism** on the GraphSAGE encoders: disabling them collapses CoraFull and Arxiv to the textbook regime (layer-6 node cosine 0.965/0.951), whereas the MAG RGCN’s per-relation weighting already keeps nodes apart (residual-off cosine 0.586, no collapse). We therefore keep residual + LayerNorm everywhere for exactly the reason the GraphSAGE case demonstrates. Together with the FullCov objective ablation (Table 1), this shows the learned component is a well-posed, depth-justified localizer, not a fragile black box.




Figure 4: Memory is the primary benefit; cold-load latency is competitive and the tradeoff is tunable. (a) The conservative streaming serve is 4.2× smaller on MAG (2.4 vs 10.2 GB); the optimized cache sweep lowers its eight-partition point to 1.9 GB (5.4×). (b) Cold-load latency, Jigsaw vs. classical CFL/DP-iso/GQL: 2.35/3.0/1.2× faster on CoraFull/Arxiv/MAG (†MAG resident enumeration is ~1 ms after loading). (c) MAG memory↔latency frontier as the cache grows (8→256): more memory buys down p50 (5.9→4.8 s) while the mean stays ≈17.6–17.9 s because of a hard-query I/O tail.




Figure 5: The six-layer encoder does not oversmooth. Per-layer partition-summary discriminability (the quantity retrieval uses) on the deployed weights: it is preserved on CoraFull, rises with depth on Arxiv and MAG, and only collapses when residual connections are disabled on the GraphSAGE encoders – isolating the residual as the anti-collapse mechanism.

5.5 Candidate-Shrinkage Cascade and Boundary Overlap

A natural objection is that one-hop overlap might expand to most of the graph. We test the per-query cascade

full graph → overlap → signature → label-prune → components on 14,400 positive method/model-query evaluation records over the 1,800-query MAG positive grid. Three facts emerge.

(1) **Overlap is broad, but ranking still removes most of the graph.** The median selected-overlap union is 59% of MAG for Jigsaw, yet after identical pruning the candidate is a median 84.5K nodes versus 544K for FilterAll, a 6.4× median reduction (this per-query cascade diagnostic uses the generic-mix encoder over 14,400 records and reports medians; the deployed-encoder means are 138K vs 444K in Table 2).

(2) **Pruning is coverage-lossless.** Recall loss occurs at retrieval/overlap, not attribute pruning:

PROPOSITION 5.1 (LOSSLESS PRUNING). *Let the pruning token $\sigma(\cdot)$ be a node attribute the matching semantics preserve, i.e. any valid embedding ϕ of Q into G satisfies $\sigma(\phi(q)) = \sigma(q)$ for every query*

node q (vertex-labelled matching, σ derived from node type and feature labels). Prune G_C to $G'_C = \{v \in G_C : \sigma(v) \in \sigma(V_Q)\}$. Then for any embedding ϕ with $\phi(V_Q) \subseteq G_C$, we have $\phi(V_Q) \subseteq G'_C$; pruning preserves node coverage and FullCov is unchanged.

Proof. For $q \in V_Q$, $\sigma(\phi(q)) = \sigma(q) \in \sigma(V_Q)$, so $\phi(q)$ meets the keep condition and is retained. Thus every matched node survives. □

The cascade signatures and exact-label filters are pure functions of each node’s stored attribute vector—node type, hashed feature bits, and (on MAG) quantized per-relation-degree features carried verbatim in the query payload, never the query’s structural degree recomputed on Q —so they are preserved under embedding and the proposition applies. Empirically, planted-match survival after overlap equals survival after pruning for every policy (Jigsaw 87% both), so the solve ceiling is set by whether the retrieved-plus-overlapped region contains the match. (3) **Candidate construction dominates latency:** mean build is 3.4 s for Jigsaw and 5.5 s for FilterAll versus 0.25 s of solving.

Boundary overlap is an essential recall operator. Removing boundary overlap cuts solve rate from 88.9% to 44.4% on MAG and from 94.4% to 86.1% on Arxiv. Its benefit is strongest for compact local structure: boundary expansion recovers missed true partitions for 94% of K-hop queries but 62% of random-walk queries, whose matches more often extend into the interior of an unselected partition. With coverage-lossless pruning, this locates the remaining errors before verification, when retrieval plus boundary overlap does not cover the match.

5.6 Query Families and Operator Necessity

The family view (Table 7) explains the aggregate rates: single and multi-fine saturate, local K-hop families remain high, and random-walk/multi-coarse are hardest across datasets. Random walks are the main stress test; FilterAll solves them at 96.0%, and 96.6% of Jigsaw’s random-walk misses occur before solving, locating the opportunity in retrieval and overlap coverage rather than in Glasgow. The held-out deployed encoder lifts random-walk 60.0%→71.0%

Table 7: Jigsaw query-family breakdown. Positive columns report exact-solve rate; negative columns report correct no-match rate. All datasets use the connected-positive protocol at the matched half-partition budget ($K=10/100/1000$); MAG rows use the deployed encoder (overall 88.6%).

Dataset	Single	K-hop	Deg-K	Walk	Multi-fine	Multi-coarse	Neg-label	Neg-struct
CoraFull	100.0	98.3	97.3	75.3	100.0	94.3	100.0	100.0
Arxiv	100.0	97.3	98.3	82.3	100.0	78.7	100.0	100.0
MAG	98.7	93.3	92.7	71.0	100.0	76.0	100.0	99.0

and multi-coarse 71.7%→76.0% over a baseline encoder, showing that learning improves the families in which spatial coverage matters most. Cumulative solved-by-budget curves rise monotonically to 94.2/92.8/88.6; on MAG the solve rate rises from 37.6% at small budgets to 88.6% at $K=1000$, confirming that candidate coverage, not verifier soundness, controls the budget frontier.

What the operator ablations establish. The cascade needs multiple operators. At matched budgets, removing overlap lowers solve rate from 88.9 to 44.4% (MAG) and 94.4 to 86.1% (Arxiv); removing signatures leaves Arxiv solve rate unchanged but expands mean pruned candidates $\sim 37\times$ ($76\rightarrow 2,815$); removing exact-label pruning and component decomposition drops MAG to 52.8% and 63.9% respectively. Diagnostic variants are kept separate from the canonical validation-selected matrix.

5.7 Negative Queries and Soundness

Jigsaw never accepts from the neural model alone; Glasgow is the only accepting step. Audited no-match accuracy is 100.0% for Cora-Full/Arxiv and 99.3–99.8% across MAG policies, with at most one false positive per method; the remaining errors are timeout/unknown outcomes. The rare false positives are benchmark-level, not verifier hallucinations: a returned mapping is a true isomorphism inside the candidate, but a perturbed “negative” may still match elsewhere. The reported metric is audited no-match accuracy under this synthetic protocol; global non-existence would require exhaustive full-graph certification, which Jigsaw does not claim.

6 RELATED WORK

Exact matchers and in-memory systems. Subgraph isomorphism is NP-complete [9]. Backtracking matchers from Ullmann [10] and VF2/VF3 [5, 25] through RI [11], LAD [12], TurboISO [18], CFL-Match [19], DP-iso [20], GraphQL [21], and the more recent VEQ [22], RapidMatch [23], and CECI [24] improve ordering, pruning, and candidate indexing, and constraint programming yields the strong Glasgow solver [1] we adopt as the accepting authority. In-depth studies [27, 28] show these systems are competitive but assume the data graph and its candidate domains are memory-resident – exactly the assumption that fails at million-node scale and that Jigsaw removes. Worst-case-optimal joins [29] and GPU matchers [26] attack the join cost rather than residence.

Scale, low memory, and distribution. Trinity [32] matches on billion-node graphs with a distributed memory cloud; DualSim [31] enumerates on a single machine by paging from disk; low-memory dataflows [30] and G-thinker [33] shard or stream the whole graph across a cluster. These stream or distribute *the entire graph*; Jigsaw

instead pays a one-time partition/index cost and bounds residence by learned relevance, exposing a complementary single-machine memory frontier.

Learned matching and indexing. Most neural approaches predict matches or similarity and let the network decide acceptance: Neuro-Match [14], GLSearch [16], subgraph counting [17], SimGNN [34], GMN [35], partition-based hierarchical matching [36], and the ISONET line [37–40]; these typically target smaller or contained queries and trade exactness for a learned score. A second group keeps exactness and uses learning only to prune: GNN-PE [15] builds path-dominance embeddings that filter candidates with no false dismissals and then verify exactly, at path granularity and on graphs larger than MAG, and is the closest system to ours. Jigsaw shares that filter-then-verify discipline – the GNN chooses *where* to verify and Glasgow alone accepts, so exactness is independent of model quality – but differs in three ways: it retrieves at partition/region granularity rather than per path, it targets *memory-bounded* serving with a tunable memory–latency frontier that GNN-PE and the resident matchers do not expose, and it trades a global no-false-dismissal guarantee for a budget-controlled completeness dial. We run the public GNN-PE release as a transfer diagnostic (Appendix B); a matched-scale head-to-head is future work. Classical structure indices GraphGrep [41] and gIndex [42] inspire our FeatureIndex baseline, which we show wins precisely when labels are near-unique. Jigsaw builds on standard components – GraphSAGE [3]/RGCN [4] encoders, METIS partitioning [7], FAISS retrieval [6], contrastive learning [43], CVaR [44], and reciprocal-rank fusion [45] – and composes them into a retrieval-constrained exact matcher.

7 DISCUSSION AND LIMITATIONS

Jigsaw is an exact-verification system with a learned search boundary. Retrieval coverage controls completeness while Glasgow preserves sound acceptance inside every candidate; FullCov makes that boundary measurable and trainable. The design targets static attributed graphs and repeated query workloads where partitioning, embeddings, and FAISS indexes amortize; a classical index leads under near-unique labels and learned retrieval becomes the better localizer as labels coarsen, which is why Jigsaw ships both behind an offline selectivity router.

Scope and non-goals. We compare against feature/structure index rankers, no-retrieval Glasgow, exhaustive FilterAll, uninformed retrievers, and scalable CFL-Match/DP-iso/GQL controls; these locate Jigsaw’s advantage in bounded residence and low-selectivity

candidate construction. We do not yet benchmark the most recent in-memory matchers (VEQ, RapidMatch, CECI); like the resident controls we do run, they assume whole-graph residence, so a matched head-to-head at million-node scale is the natural next comparison rather than a threat to the residence argument. Jigsaw does *not* certify global non-existence – a bounded no-match is not a proof no match exists in G – and the current benchmark uses planted connected subgraphs on static, single-topology citation graphs, so multi-topology and dynamic-graph generalization remain open. The walk family is the standing recall bottleneck, and the inference-time probes locate its residual errors in representation-level footprint coverage, giving a concrete next learning target. Out-of-core and distributed matchers remain complementary: combining Jigsaw’s learned residence bound with sharded verification is promising future work.

8 CONCLUSION

Jigsaw turns learned graph retrieval into an execution mechanism for exact search: the model predicts which regions deserve memory and verification, while Glasgow remains the sole accepting authority. On full MAG this changes the tested operating point from 0/15 direct solves on the K -hop control to 88.6% exact recovery over 1,800 planted positives, while streaming at 2.4 GB instead of 10.2 GB for whole-graph residence. FullCov aligns learning with the decisive failure event – missing even one required partition – while boundary overlap and coverage-lossless pruning make the verifier usable under bounded memory. Under near-unique labels a resident classical index dominates; under weak labels whole-graph candidate construction fails and Jigsaw’s bounded candidate supplies the systems advantage, so Jigsaw is retriever-agnostic by construction. The result is exact subgraph verification made usable on million-node graphs without replacing exactness by a neural guess.

ACKNOWLEDGMENTS

Generative AI tools assisted with implementation, experiment orchestration, data-validation scripts, and manuscript language editing. The author reviewed and verified all AI-assisted code, analyses, numerical claims, and text, and takes full responsibility for the final work.

A ENCODER AND TRAINING DETAILS

Beyond the architecture of Section 3.1 (Figure 2), the readout is an MLP plus a skip projection before L_2 normalization. The message-passing block is GraphSAGE [3] on CoraFull/Arxiv – a matched objective ablation finds it at least as strong as GIN [2] at the budgets the cascade uses – and a relation-aware six-layer RGCN [4] on MAG (per-relation bases over $2R$ symmetrized relations); non-paper MAG nodes, which lack features, receive node-type one-hot and per-relation-degree features. Training combines three signals: paired fine/coarse InfoNCE, hard negatives, and the FullCov coverage term; positive partitions are live re-encoded so coverage gradients update required positives rather than a stale bank. The MAG checkpoint is selected by validation FullCov (seeds 31415/27182), disjoint from the benchmark seeds. Queries are inductive with respect to an indexed target and need no fine-tuning; a new target pays a

Table 8: Representative deployment costs. Hierarchy loading and combined embedding-plus-index construction are measured setup runs; training is a one-time dataset cost; rank time is query encoding plus FAISS search and excludes candidate construction.

Target	Hier. load	Embed + index	Enc. training	Rank/query
Arxiv (200 coarse)	3.5 s	1.9 s	2–3 h (9,000 steps)	5.1 ms
MAG (2,000 coarse)	110.9 s	13.7 s	~3 h (5,000 steps)	21.9 ms

Table 9: Operator-necessity ablation at matched half-partition budgets. The shared diagnostic has 48 queries per dataset; solve/candidate statistics use its 36 planted positives (MAG at $K=1,000/2,000$, Arxiv at $K=100/200$).

Variant	OGBN-MAG		OGBN-Arxiv	
	Solve %	Cand.	Solve %	Cand.
Full Jigsaw	88.9	150.6K	94.4	76
No overlap	44.4	152.5K	86.1	88
No signature	88.9	150.7K	94.4	2,815
No components	63.9	143.0K	94.4	76
No exact-label	52.8	223.6K	94.4	244

one-time METIS hierarchy, partition-embedding, and FAISS build (Table 8).

B GNN-PE PUBLIC-RELEASE DIAGNOSTIC

Beyond the direct-solver control (Table 3, 0/15 on Arxiv/MAG), we ran the public GNN-PE [15] release. On CoraFull the released pipeline passes its self-test but, after class-label export and build repairs, recovers 7/24 planted positives at edit distance $e=4$ (0/6 at size 100). The release exposes an inference path but no supervised training recipe for this protocol, so the number measures public-release transfer, not a retrained reproduction; we report it as a transparency control rather than a head-to-head result.

C OPERATOR-NECESSITY ABLATION

Table 9 gives the per-operator breakdown summarized in Section 5.6. The equal-budget comparison removes the full-budget confound. On MAG, removing one-hop overlap halves solve rate (88.9→44.4%) while signatures are inert; on Arxiv, signatures are the main candidate-control mechanism (mean pruned domain expands $\sim 37\times$, 76→2,815, when removed) and overlap has a smaller but real effect (94.4→86.1%). Exact-label pruning and component decomposition remain important on MAG (88.9→52.8% and 88.9→63.9%).

D PARTITION AND OVERLAP SCALE

For a selected coarse partition, the overlap operator adds only boundary nodes from partitions sharing an edge with it, not the bordering partitions in full. On MAG a median partition (969 nodes) borders 948 of 2,000 coarse partitions but admits only $\approx 4,252$ boundary nodes, so the expanded region is $\approx 5,228$ nodes – a $5.4\times$ expansion, not the $\approx 921K$ a whole-neighbor union would give. That loose

Table 10: Inference-time remedies on the same 1,800-query MAG positive workload: FullCov@1000 (%). Each displayed spatial family has $n=300$ queries.

Method	k-hop	degree-k-hop	random-walk
Single-vector (current)	14.7	17.3	9.0
Subquery MaxSim (query-only)	14.3	17.7	8.0
Finer partition parent	12.7	16.7	8.3
Fine-level overlap (1-hop)	13.3	17.3	10.0
Fine-level overlap (2-hop)	13.3	18.3	9.3
Dual-granularity late interaction	12.7	16.0	8.7
Coarse diffusion	17.0	19.7	10.3

bound is a reachability figure the operator never realizes; reporting it as “the overlap” overstates MAG breadth by more than $170\times$. Candidate breadth arises only when the cascade unions boundary overlap across hundreds of retrieved partitions, and because many boundary nodes are shared the union saturates between 59% (Jigsaw) and 83% (uninformed policies) of MAG at the largest budgets.

E RETRIEVAL-REMEDY FORECLOSURE AND WITHIN-MAG SELECTIVITY

A natural question is whether a cheaper inference-time change to retrieval would close the spatial-family gap without retraining. It does not. On MAG, subquery multi-vector MaxSim, finer partition parents, one- and two-hop fine-level overlap, dual-granularity late interaction, and coarse diffusion all leave FullCov@1000 within ≈ 2 points of the single-vector baseline on the k-hop/degree-k-hop/random-walk families (Table 10); the best gain is $\approx +2.3$ (coarse diffusion) and a coarse neighbor-stitch expansion *reduces* coverage by 8–13 points. The cause is density: the partition graph is high-degree at every granularity (median coarse degree 948; mean fine degree 568 of 10,000), so neighborhood-following broadens the candidate rather than tracing the query footprint. This directs further improvement toward coverage-aware representation learning rather than inference-time patching.

The selectivity crossover of Section 5.2 also appears *within* MAG. Holding the 1.9M-node target fixed and shrinking the md5 codebook from 10^6 to 10^2 labels drives the median partitions spanned by a label from 1 to 1,938; the learned ranker never uses this codebook, so its coverage is flat, while FeatureIndex begins at the ceiling and falls monotonically, crossing below the learned ranker family by family (Table 11). This reproduces the cross-dataset inversion on a single deployed target, confirming the learned advantage is a low-selectivity phenomenon rather than a dataset artifact.

REFERENCES

- [1] McCreesh, C., Prosser, P., Trimble, J.: The Glasgow Subgraph Solver: Using Constraint Programming to Tackle Hard Subgraph Isomorphism Problem Variants. In: CP 2020, LNCS, vol. 12333, pp. 316–333. Springer (2020)
- [2] Xu, K., Hu, W., Leskovec, J., Jegelka, S.: How Powerful are Graph Neural Networks? In: ICLR (2019)
- [3] Hamilton, W.L., Ying, R., Leskovec, J.: Inductive Representation Learning on Large Graphs. In: NeurIPS, pp. 1024–1034 (2017)
- [4] Schlichtkrull, M., Kipf, T.N., Bloem, P., van den Berg, R., Titov, I., Welling, M.: Modeling Relational Data with Graph Convolutional Networks. In: ESWC, pp. 593–607 (2018)
- [5] Cordella, L.P., Foggia, P., Sansone, C., Vento, M.: A (Sub)Graph Isomorphism Algorithm for Matching Large Graphs. IEEE TPAMI 26(10), 1367–1372 (2004)

Table 11: Within-MAG label-selectivity sweep (FullCov@1000, feature-driven families, $n=300$ each). As the md5 codebook shrinks, partitions/label grows and FeatureIndex coverage falls below the *label-independent* learned ranker (bottom row, flat by construction).

Median parts/label	Single	Multi-fine	Multi-coarse
1 (near-unique)	100.0	100.0	99.0
7	100.0	100.0	99.0
74	100.0	100.0	94.0
636	100.0	100.0	59.3
1,938	76.0	78.7	14.0
Learned (flat)	99.3	99.7	65.0

- [6] Johnson, J., Douze, M., Jégou, H.: Billion-Scale Similarity Search with GPUs. IEEE Trans. Big Data 7(2), 535–547 (2021)
- [7] Karypis, G., Kumar, V.: A Fast and High Quality Multilevel Scheme for Partitioning Irregular Graphs. SIAM J. Sci. Comput. 20(1), 359–392 (1998)
- [8] Hu, W., et al.: Open Graph Benchmark: Datasets for Machine Learning on Graphs. In: NeurIPS (2020)
- [9] Garey, M.R., Johnson, D.S.: Computers and Intractability: A Guide to the Theory of NP-Completeness. W.H. Freeman (1979)
- [10] Ullmann, J.R.: An Algorithm for Subgraph Isomorphism. J. ACM 23(1), 31–42 (1976)
- [11] Bonnici, V., Giugno, R., Pulvirenti, A., Shasha, D., Ferro, A.: A Subgraph Isomorphism Algorithm and Its Application to Biochemical Data. BMC Bioinformatics 14(Suppl 7), S13 (2013)
- [12] Solnon, C.: AllDifferent-Based Filtering for Subgraph Isomorphism. Artificial Intelligence 174(12–13), 850–864 (2010)
- [13] Bojchevski, A., Günnemann, S.: Deep Gaussian Embedding of Graphs: Unsupervised Inductive Learning via Ranking. In: ICLR (2018)
- [14] Ying, R., et al.: Neural Subgraph Matching. arXiv:2007.03092 (2020)
- [15] Ye, Y., Lian, X., Chen, M.: Efficient Exact Subgraph Matching via GNN-based Path Dominance Embedding. Proc. VLDB Endow. 17(7), 1628–1641 (2024)
- [16] Liu, Y., et al.: GLSearch: Maximum Common Subgraph Detection via Learning to Search. In: ICML (2021)
- [17] Liu, X., Pan, H., He, M., Song, Y., Jiang, X., Shang, L.: Neural Subgraph Isomorphism Counting. In: KDD (2020)
- [18] Han, W.S., Lee, J., Lee, J.H.: TurboISO: Towards Ultrafast and Robust Subgraph Isomorphism Search in Large Graph Databases. In: SIGMOD, pp. 337–348 (2013)
- [19] Bi, F., Chang, L., Lin, X., Qin, L., Zhang, W.: Efficient Subgraph Matching by Postponing Cartesian Products. In: SIGMOD, pp. 1199–1214 (2016)
- [20] Han, M., Kim, H., Gu, G., Park, K., Han, W.S.: Efficient Subgraph Matching: Harmonizing Dynamic Programming, Adaptive Matching Order, and Failing Set Together. In: SIGMOD, pp. 1429–1446 (2019)
- [21] He, H., Singh, A.K.: Graphs-at-a-Time: Query Language and Access Methods for Graph Databases. In: SIGMOD, pp. 405–418 (2008)
- [22] Kim, H., Choi, Y., Park, K., Lin, X., Hong, S.H., Han, W.S.: Versatile Equivalences: Speeding up Subgraph Query Processing and Subgraph Matching. In: SIGMOD, pp. 925–937 (2021)
- [23] Sun, S., Sun, X., Che, Y., Luo, Q., He, B.: RapidMatch: A Holistic Approach to Subgraph Query Processing. Proc. VLDB Endow. 14(2), 176–188 (2020)
- [24] Bhattarai, B., Liu, H., Huang, H.H.: CECI: Compact Embedding Cluster Index for Scalable Subgraph Matching. In: SIGMOD, pp. 1447–1462 (2019)
- [25] Carletti, V., Foggia, P., Saggese, A., Vento, M.: Challenging the Time Complexity of Exact Subgraph Isomorphism for Huge and Dense Graphs with VF3. IEEE Trans. Pattern Anal. Mach. Intell. 40(4), 804–818 (2018)
- [26] Zeng, L., Zou, L., Özsu, M.T., Hu, L., Zhang, F.: GSI: GPU-friendly Subgraph Isomorphism. In: ICDE, pp. 1249–1260 (2020)
- [27] Sun, S., Luo, Q.: In-Memory Subgraph Matching: An In-depth Study. In: SIGMOD, pp. 1083–1098 (2020)
- [28] Zhang, Z., Lu, Y., Zheng, W., Lin, X.: A Comprehensive Survey and Experimental Study of Subgraph Matching: Trends, Unbiasedness, and Interaction. Proc. ACM Manag. Data 2(1), Article 60 (2024)
- [29] Mhedhbi, A., Salihoglu, S.: Optimizing Subgraph Queries by Combining Binary and Worst-Case Optimal Joins. Proc. VLDB Endow. 12(11), 1692–1704 (2019)
- [30] Ammar, K., McSherry, F., Salihoglu, S., Joglekar, M.: Distributed Evaluation of Subgraph Queries Using Worst-Case Optimal Low-Memory Dataflows. Proc. VLDB Endow. 11(6), 691–704 (2018)

- [31] Kim, H., Lee, J., Bhowmick, S.S., Han, W.-S., Lee, J., Ko, S., Jarrah, M.H.A.: DUAL-SIM: Parallel Subgraph Enumeration in a Massive Graph on a Single Machine. In: SIGMOD, pp. 1231–1245 (2016)
- [32] Sun, Z., Wang, H., Wang, H., Shao, B., Li, J.: Efficient Subgraph Matching on Billion Node Graphs. *Proc. VLDB Endow.* **5**(9), 788–799 (2012)
- [33] Yan, D., Guo, G., Chowdhury, M.M.R., Özsu, M.T., Ku, W.-S., Lui, J.C.S.: G-thinker: A Distributed Framework for Finding Qualified Subgraphs in a Big Graph with Load Balancing. In: ICDE (2020)
- [34] Bai, Y., Ding, H., Bian, S., Chen, T., Sun, Y., Wang, W.: SimGNN: A Neural Network Approach to Fast Graph Similarity Computation. In: WSDM, pp. 384–392 (2019)
- [35] Li, Y., Gu, C., Dullien, T., Vinyals, O., Kohli, P.: Graph Matching Networks for Learning the Similarity of Graph Structured Objects. In: ICML, PMLR **97**, 3835–3845 (2019)
- [36] Xu, H., Duan, Z., Wang, Y., Feng, J., Chen, R., Zhang, Q., Xu, Z.: Graph Partitioning and Graph Neural Network Based Hierarchical Graph Matching for Graph Similarity Computation. *Neurocomputing* **439**, 348–362 (2021)
- [37] Roy, I., Velugoti, V.S.B.R., Chakrabarti, S., De, A.: Interpretable Neural Subgraph Matching for Graph Retrieval. In: AAAI **36**(7), 8115–8123 (2022)
- [38] Ramachandran, A., Raj, V., Roy, I., Chakrabarti, S., De, A.: Iteratively Refined Early Interaction Alignment for Subgraph Matching Based Graph Retrieval. In: *NeurIPS* **37** (2024)
- [39] Raj, V., Roy, I., Ramachandran, A., Chakrabarti, S., De, A.: Charting the Design Space of Neural Graph Representations for Subgraph Matching. In: ICLR (2025)
- [40] Ying, Y., Dai, Y., Li, W., et al.: Neural Graph Navigation for Intelligent Subgraph Matching. In: AAAI **40**(19), 16181–16189 (2026)
- [41] Shasha, D., Wang, J.T.L., Giugno, R.: Algorithmics and Applications of Tree and Graph Searching. In: PODS, pp. 39–52 (2002)
- [42] Yan, X., Yu, P.S., Han, J.: Graph Indexing: A Frequent Structure-based Approach. In: SIGMOD, pp. 335–346 (2004)
- [43] van den Oord, A., Li, Y., Vinyals, O.: Representation Learning with Contrastive Predictive Coding. *arXiv:1807.03748* (2018)
- [44] Rockafellar, R.T., Uryasev, S.: Optimization of Conditional Value-at-Risk. *Journal of Risk* **2**(3), 21–41 (2000)
- [45] Cormack, G.V., Clarke, C.L.A., Buettcher, S.: Reciprocal Rank Fusion Outperforms Condorcet and Individual Rank Learning Methods. In: SIGIR, pp. 758–759 (2009)